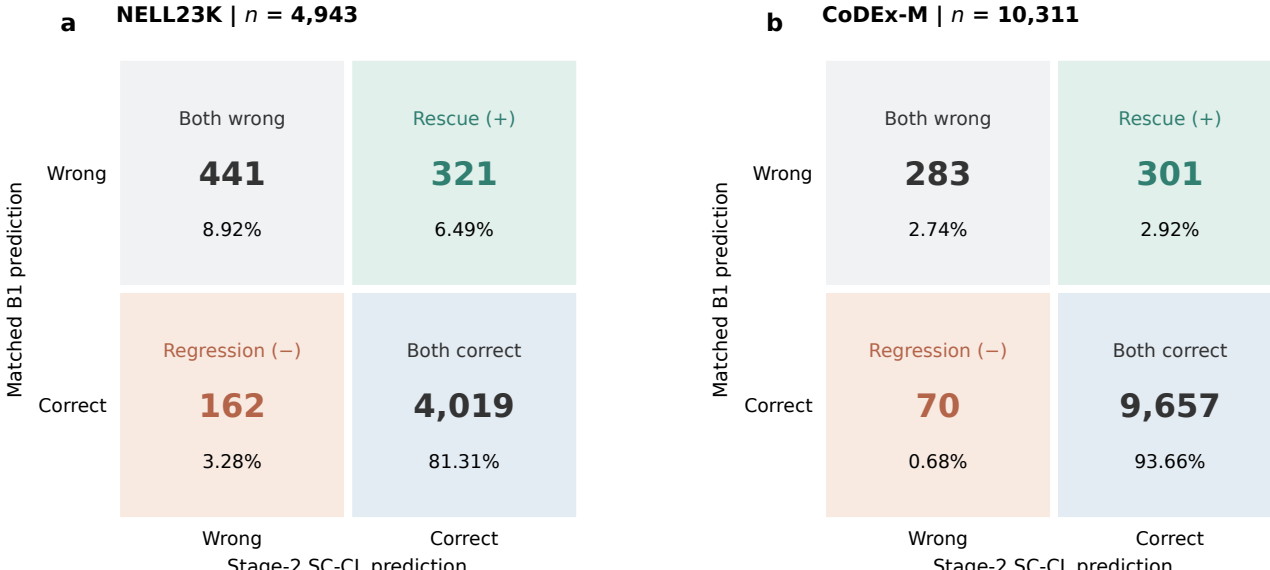

Each cell: number of test questions and percentage of that dataset's full test cohort.

Rescue − regression = net gain: 321 − 162 = 159 (NELL23K); 301 − 70 = 231 (CoDEx-M).

Single seed (2026); same prompts, candidate order, targets and question IDs within each dataset.