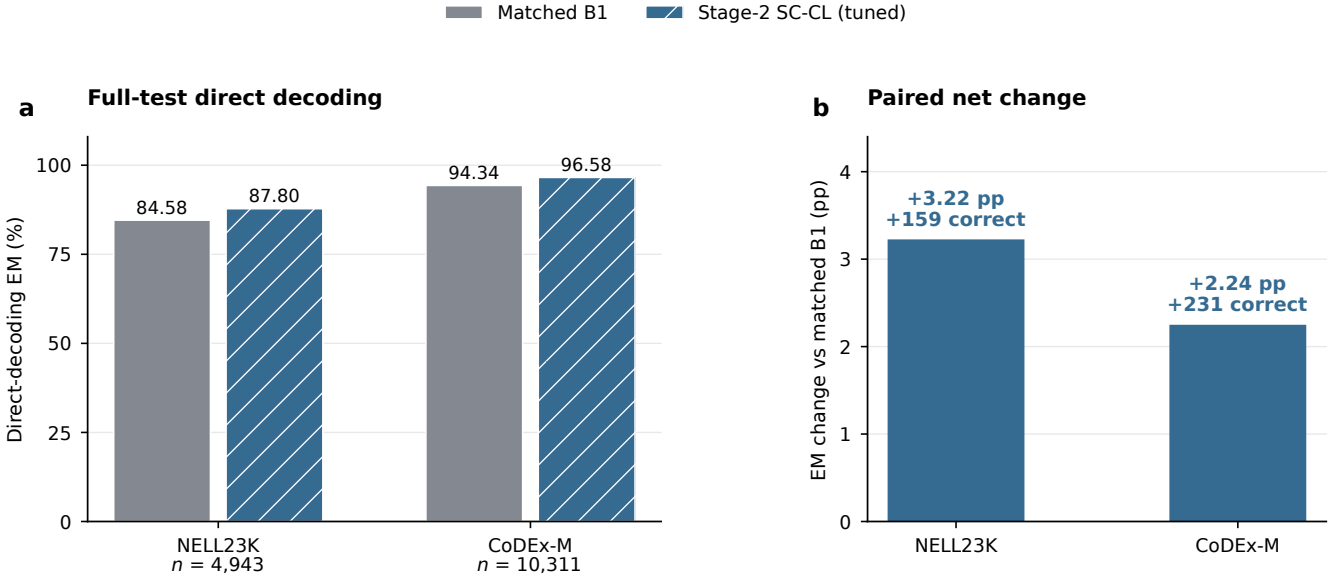

One training seed (2026); validation-selected configurations; no seed-level error bars.

Local protocol comparison; tuning and optimizer differences prevent single-component attribution.

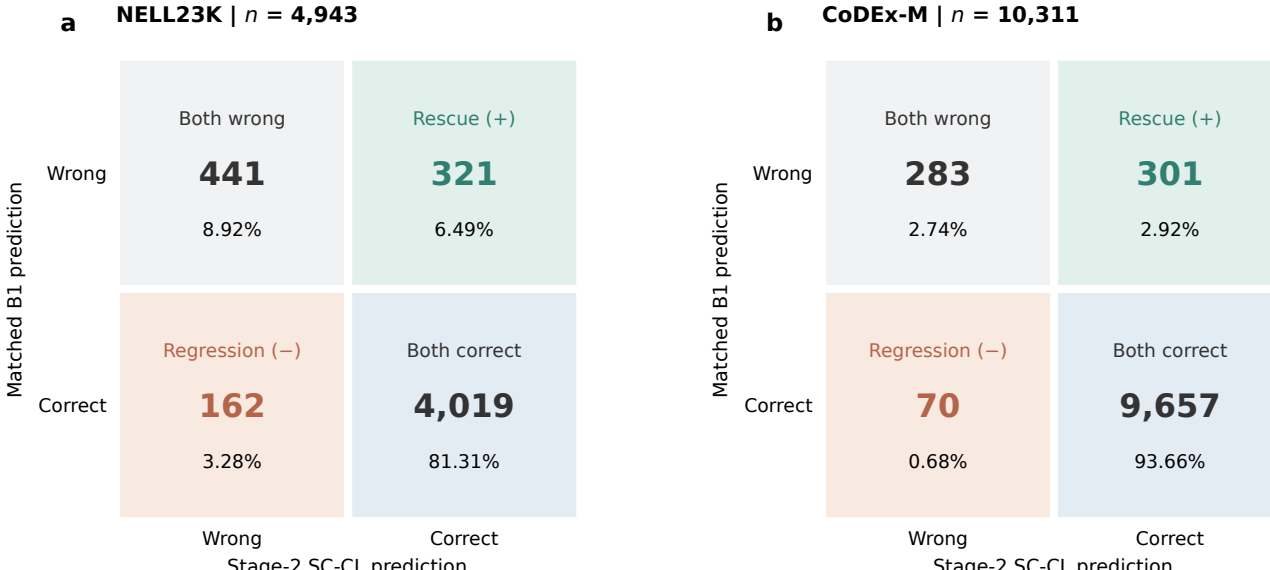

Each cell: number of test questions and percentage of that dataset's full test cohort.

Rescue − regression = net gain: 321 − 162 = 159 (NELL23K); 301 − 70 = 231 (CoDEx-M).

Single seed (2026); same prompts, candidate order, targets and question IDs within each dataset.

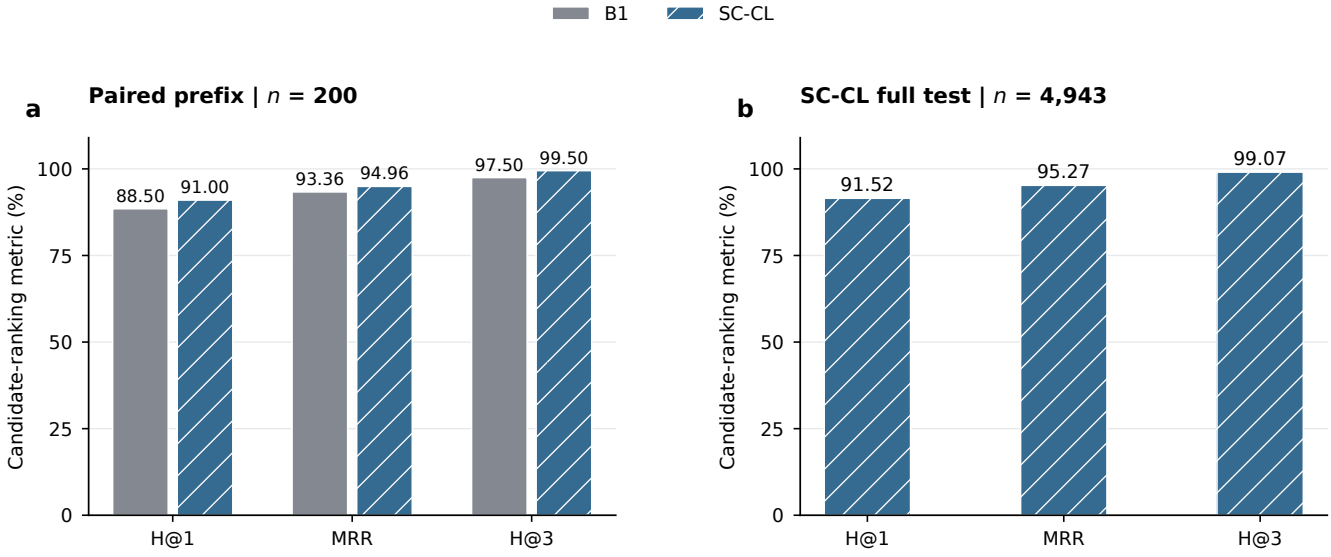

Historical NELL23K run; teacher-forced scoring of the same ten candidates; seed 2026.

Prefix follows test-file order; full-test B1 reranking is unavailable; training budgets differ.

Ranking models and baseline differ from those used for the latest direct-decoding evaluation.