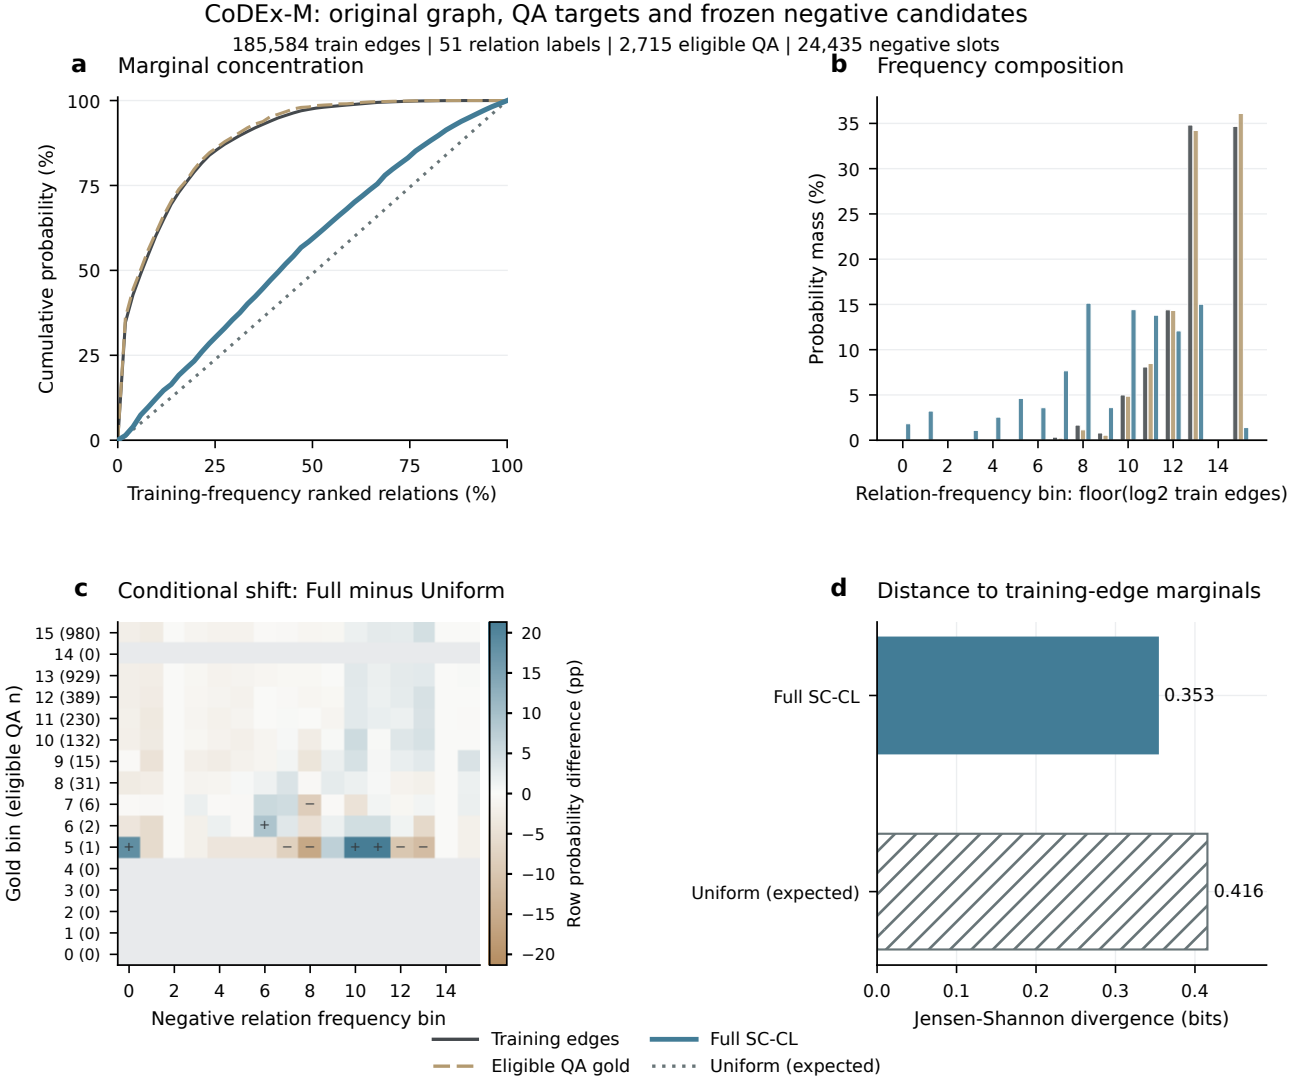
CoDEx-M: original graph, QA targets and frozen negative candidates
185,584 train edges | 51 relation labels | 2,715 eligible QA | 24,435 negative slots
a
Marginal concentration
Cumulative probability (%)
Training-frequency ranked relations (%)
b
Frequency composition
Probability mass (%)
Relation-frequency bin: floor(log2 train edges)
c
Conditional shift: Full minus Uniform
Gold bin (eligible QA n)
Negative relation frequency bin
Row probability difference (pp)
d
Distance to training-edge marginals
Full SC-CL
0.353
Uniform (expected)
0.416
Jensen-Shannon divergence (bits)
Training edges
Eligible QA gold
Full SC-CL
Uniform (expected)
CoDEx-M: no trained Uniform/quota controls. Dotted lines and hatched bars are analytic legal-pool expectations.