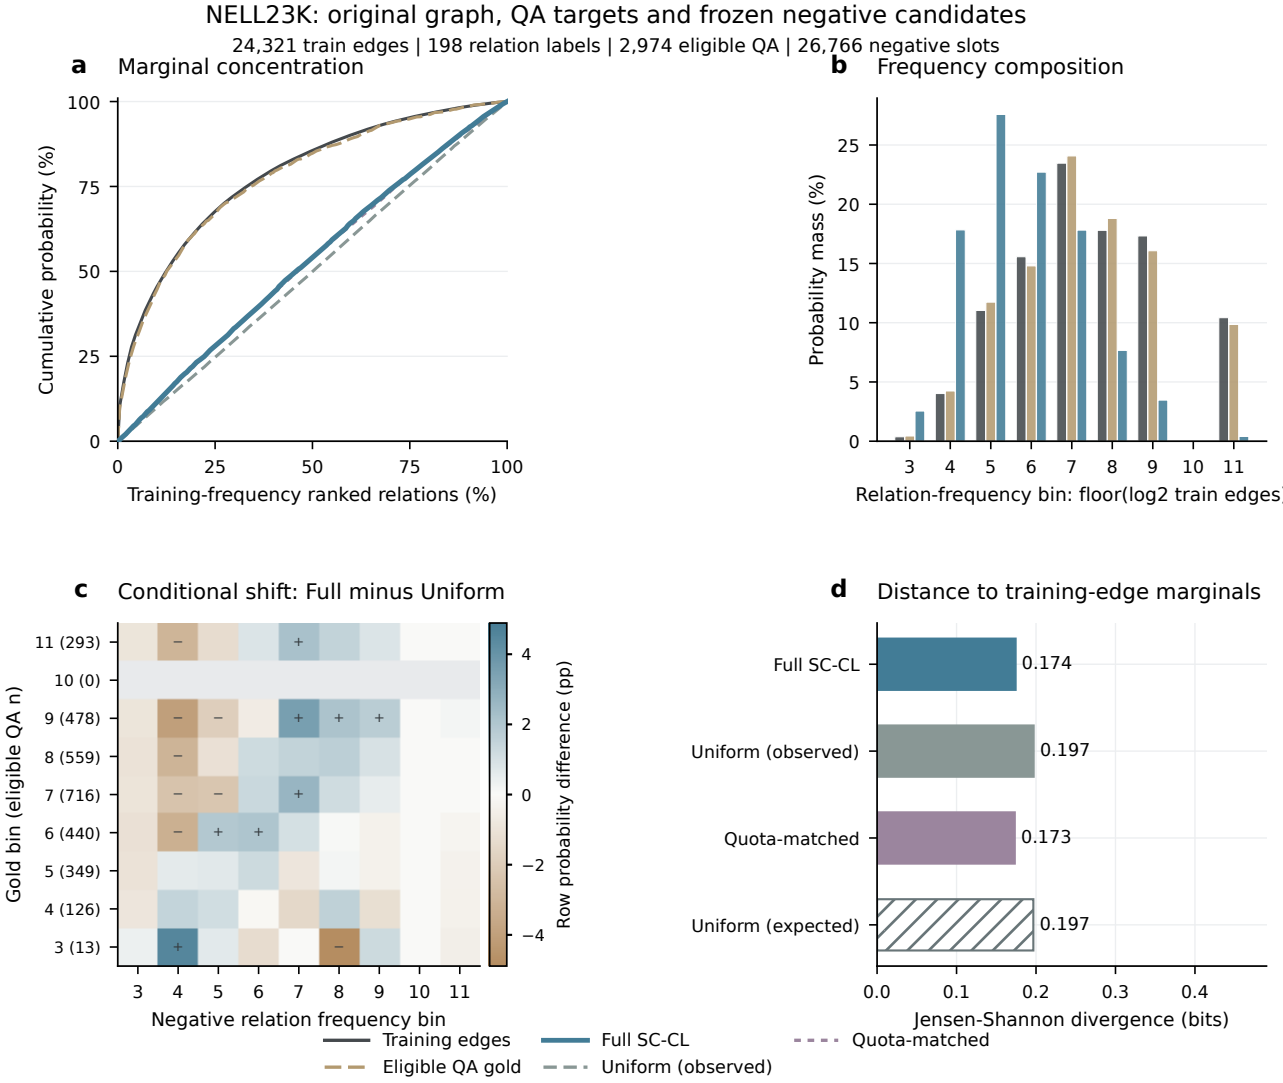
NELL23K: original graph, QA targets and frozen negative candidates
24,321 train edges | 198 relation labels | 2,974 eligible QA | 26,766 negative slots
a
Marginal concentration
Cumulative probability (%)
Training-frequency ranked relations (%)
b
Frequency composition
Probability mass (%)
Relation-frequency bin: floor(log2 train edges)
c
Conditional shift: Full minus Uniform
Gold bin (eligible QA n)
Negative relation frequency bin
Row probability difference (pp)
d
Distance to training-edge marginals
Full SC-CL
Uniform (observed)
Quota-matched
Uniform (expected)
0.174
0.197
0.173
0.197
Jensen-Shannon divergence (bits)
Training edges
Eligible QA gold
Full SC-CL
Uniform (observed)
Quota-matched
Uniform expectation uses each eligible question's legal pool; gray heatmap rows have no observed gold QA.